


# Deep Learning for Computer Vision and Scene Understanding



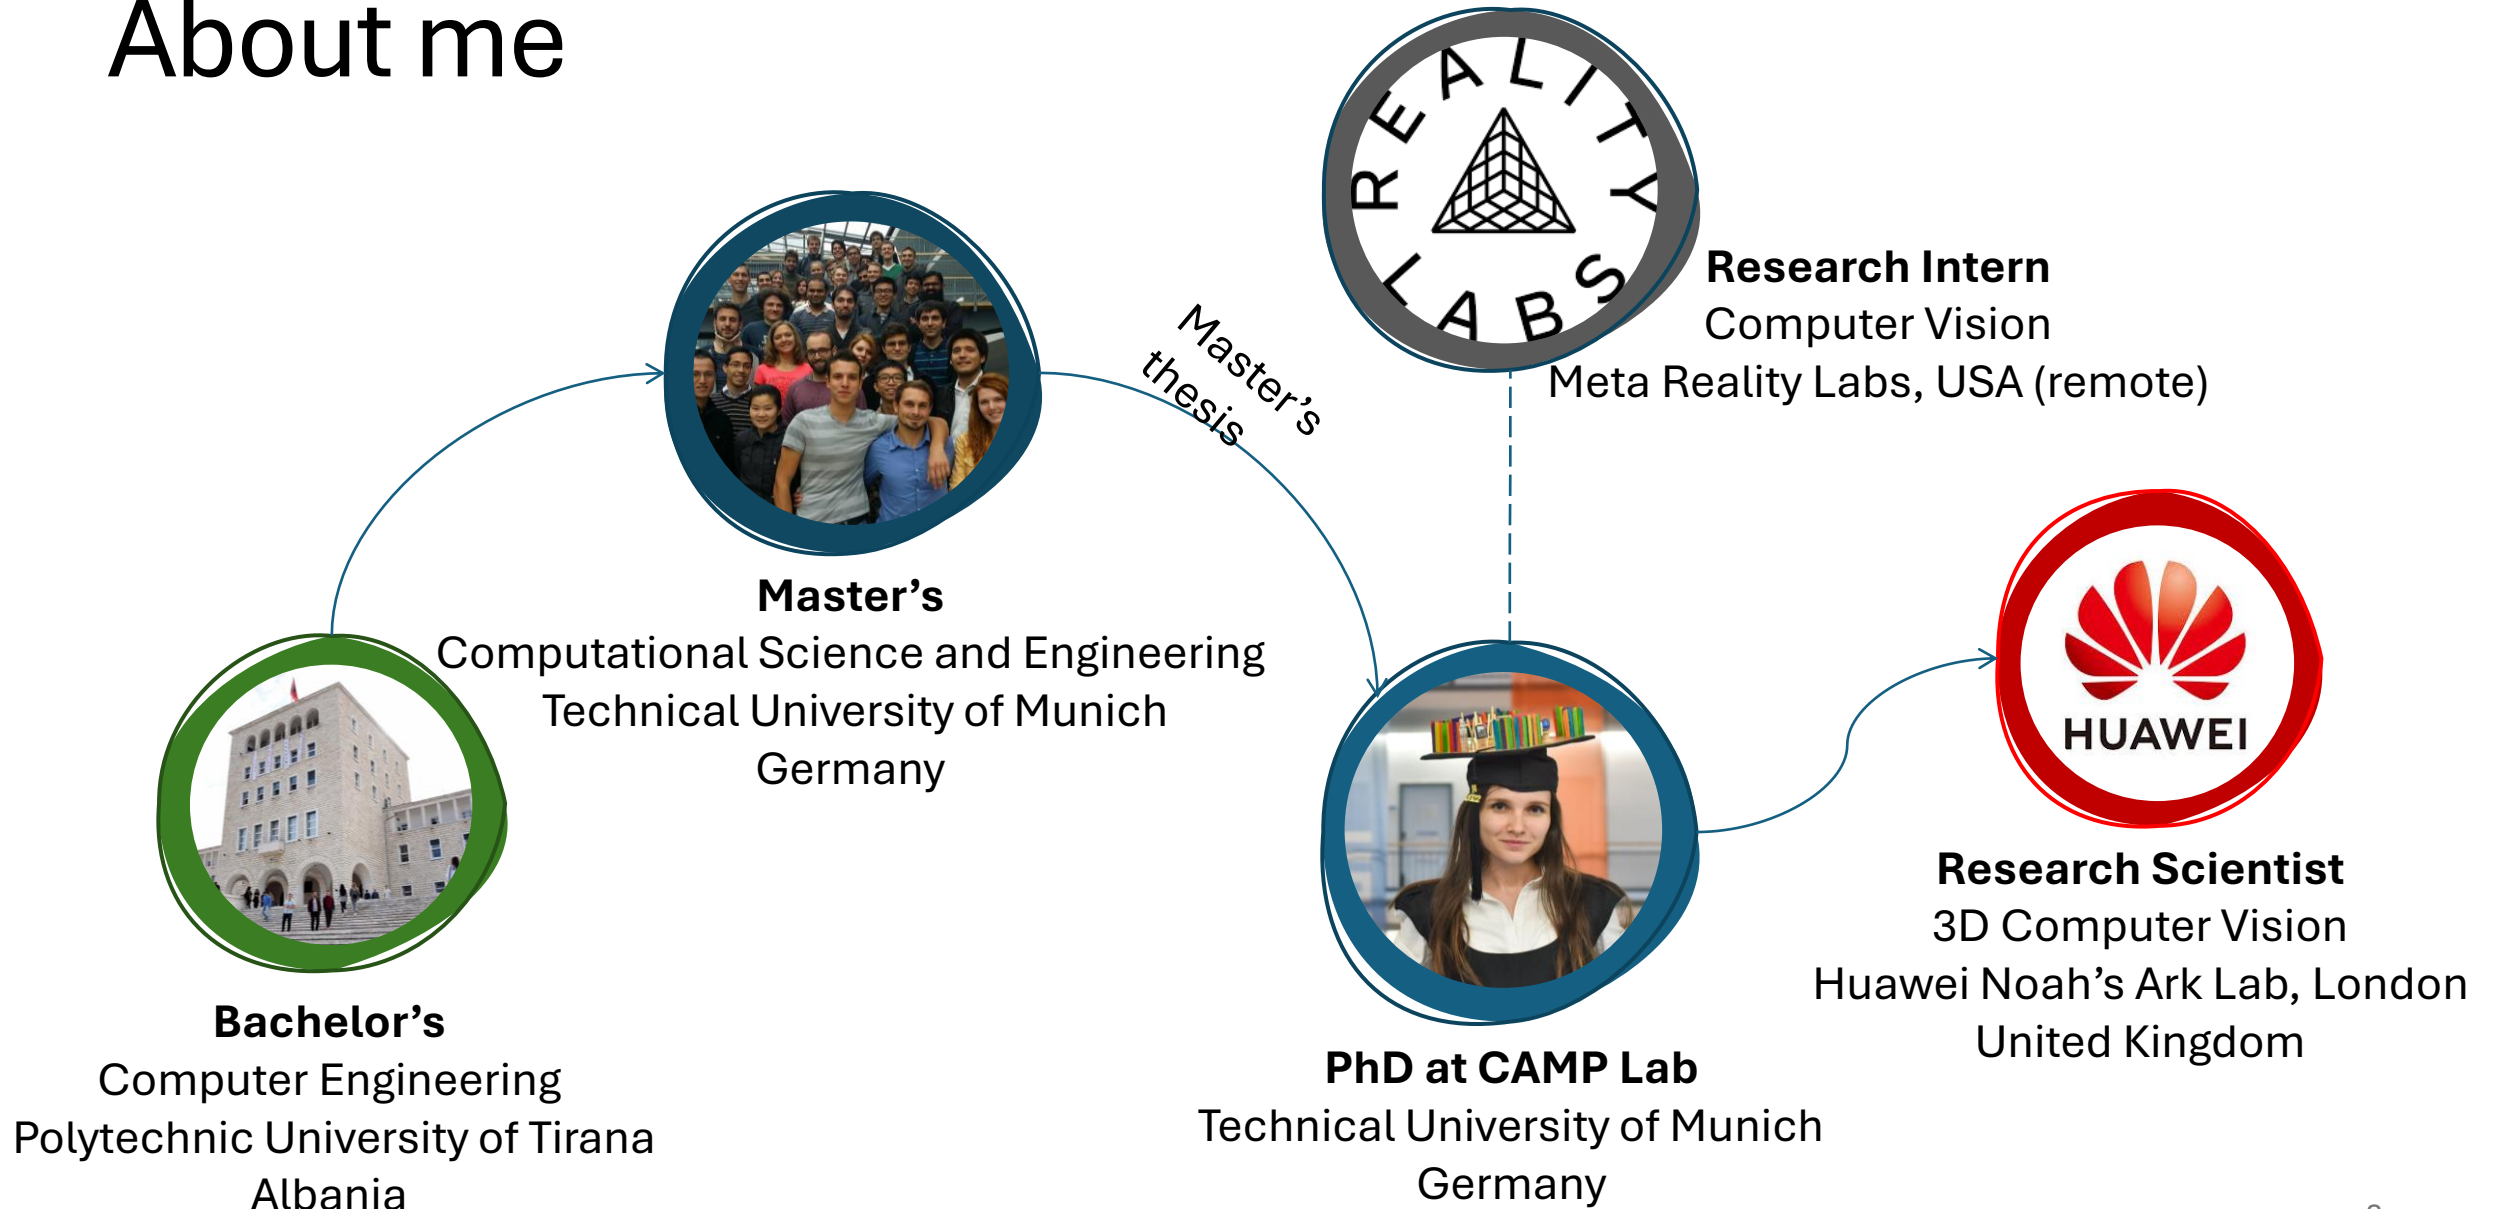
# Dr. Helisa Dhamo

## Research Scientist in 3D Vision

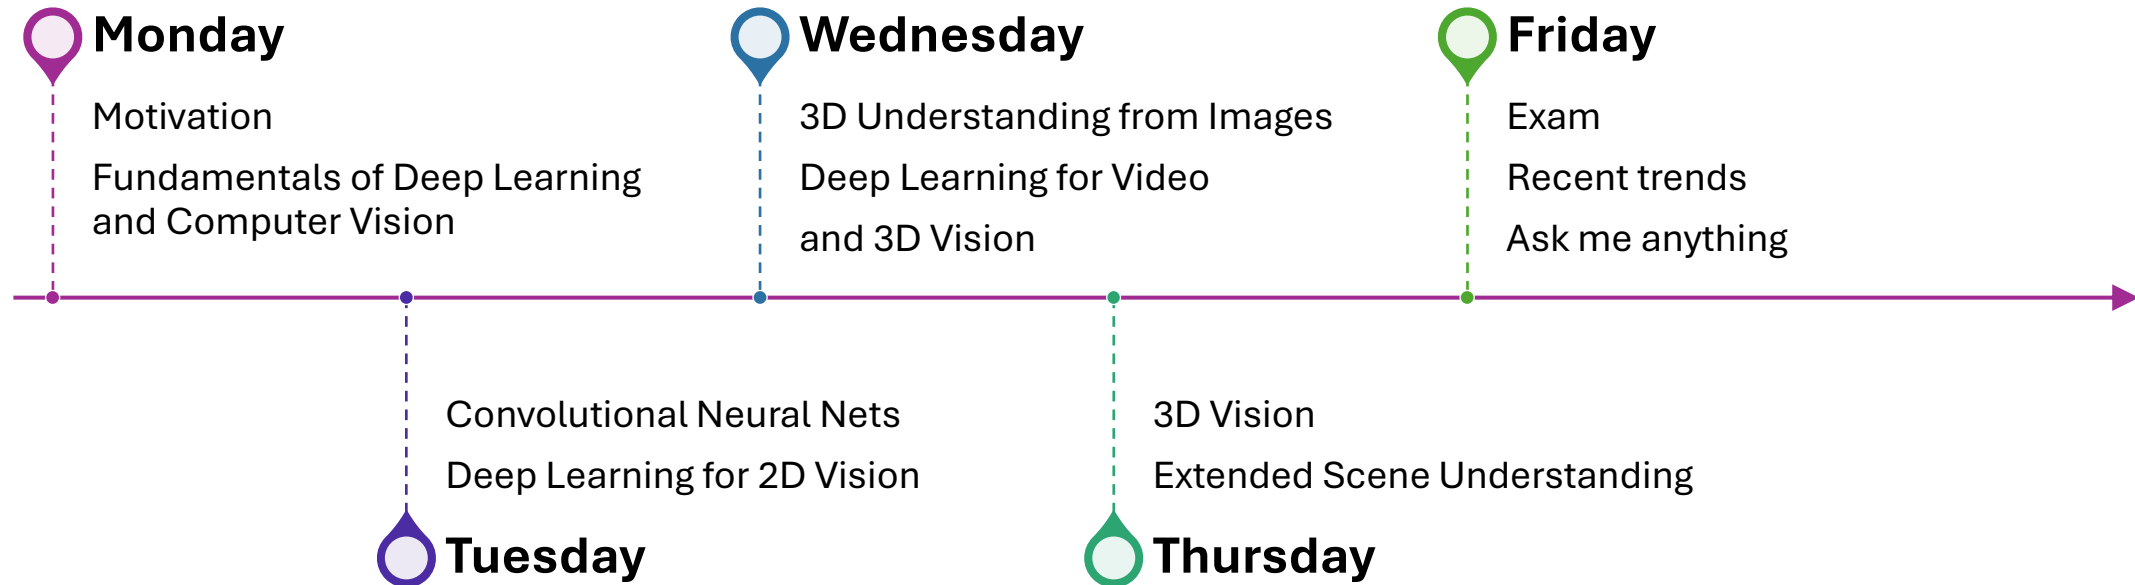
# Huawei Noah's Ark Lab London



# About me



# Course outline



# Grading

Theoretical Exam on Friday 2 August 2024

- Mostly multichoice questions
- 1-2 practical exercises

Practical coding exercises for **bonus** grade

- Send back via email before 11 August 2024





Motivation and Fundamentals

# Lecture 1



?

What does the computer see?



What do I see



27	42	62	84	89	89	0	40	66	68	89	0	0	0
99	19	89	84	89	89	0	19	89	89	89	19	89	84
89	24	53	126	249	20	9	9	0	5	44	22	8	
48													
43	10	42	4	255	250	48	43	22	27	64	86	33	
43													
0	3	88	77	32	123	43	62	84	89	89	0	0	
12													
12	2	2	5	3	60	2	70	123	43	120	115	4	
20													
66	45	52	0	126	249	20	9	9	0	123	43		
22													
50	3	9	0	5	44	22	8	97	123	43	120		
89													
88	4	43	62	84	89	89	0	40	66	68	80		
120													
0	88	77	32	123	43	120	115	40	50	88	77	3	
22													
15	0	9	0	5	44	22	8	97	88	77	32		
62													
25	88	77	32	123	43	62	84	89	0	76	4		
122													
88	3	9	0	5	44	22	8	97	156	149	120		
150													
12	12	43	120	115	40	50	88	255	250	48	43		
122													
2	6	9	0	5	44	22	8	97	88	77	32		
100													
	43	62	84	89	89	0	123	43	120	115	40		

What does the computer see?



What do I see

## ***Computer Vision***

Field of Computer Science that aims to make sense at image/video inputs, i.e. **identify, understand** and **extract relevant information**




## ***Scene Understanding***

Aspect of Computer Vision that aims to identify and analyse objects and their context (surrounding scene, relations to other objects)

27	42	62	84	89	89	0	40	66	68	89	0	0	0
99	19	89	84	89	89	0	19	89	89	89	19	89	84
89	24	53	126	249	20	9	9	0	5	44	22	8	4
48													
43	10	42	4	255	250	48	43	22	27	64	86	33	4
43													
0	3	88	77	32	123	43	62	84	89	89	0	0	0
12													
12	2	2	5	3	60	2	70	123	43	120	115	4	0
20													
66	45	52	0	126	249	20	9	9	0	123	43	4	0
22													
50	3	9	0	5	44	22	8	97	123	43	120	4	0
89													
88	4	43	62	84	89	89	0	40	66	68	80	4	0
120													
0	88	77	32	123	43	120	115	40	50	88	77	3	0
22													
15	0	9	0	5	44	22	8	97	88	77	32	4	0
62													
25	88	77	32	123	43	62	84	89	0	76	4	4	0
122													
88	3	9	0	5	44	22	8	97	156	149	120	4	0
150													
12	12	43	120	115	40	50	88	255	250	48	43	4	0
122													
2	6	9	0	5	44	22	8	97	88	77	32	4	0
100													
43	62	84	89	89	0	123	43	120	115	40	4	0	0

What does the computer see?



dog

ball

What do I see

# Computer Vision Applications



# Image segmentation








# Human pose estimation













# Object 6D pose estimation





# Depth estimation





# Image generation and editing

*"Swap sunflowers with roses"*



*"Add fireworks to the sky"*



*"Replace the fruits with cake"*



*"What would it look like if it were snowing?"*



*"Turn it into a still from a western"*




*"Make his jacket out of leather"*



# Between 2D and 3D

## Rendering (a Graphics problem)

- Given a 3D model of the scene (3D mesh, materials, lighting), and a camera, obtain an image



## Inverse Rendering (a Vision problem)



- Given an image of a scene, infer the 3D model
- Under-constrained, ill-posed







# Computer Vision: Pinhole camera model

Image formation



# Computer Vision

The pinhole camera model - Image formation





$f$  - Focal length  
 $z$  - depth

$$\frac{v}{f} = \frac{y}{z}, \quad \frac{u}{f} = \frac{x}{z}$$

# Computer Vision

The pinhole camera model - Image formation



$f$  - Focal length

$z$  - depth

$c = (c_x, c_y)$  - optical center



$$v = f \frac{y}{z}$$

$$\text{and } u = f \frac{x}{z}$$



# Computer Vision


## The pinhole camera model - Image formation



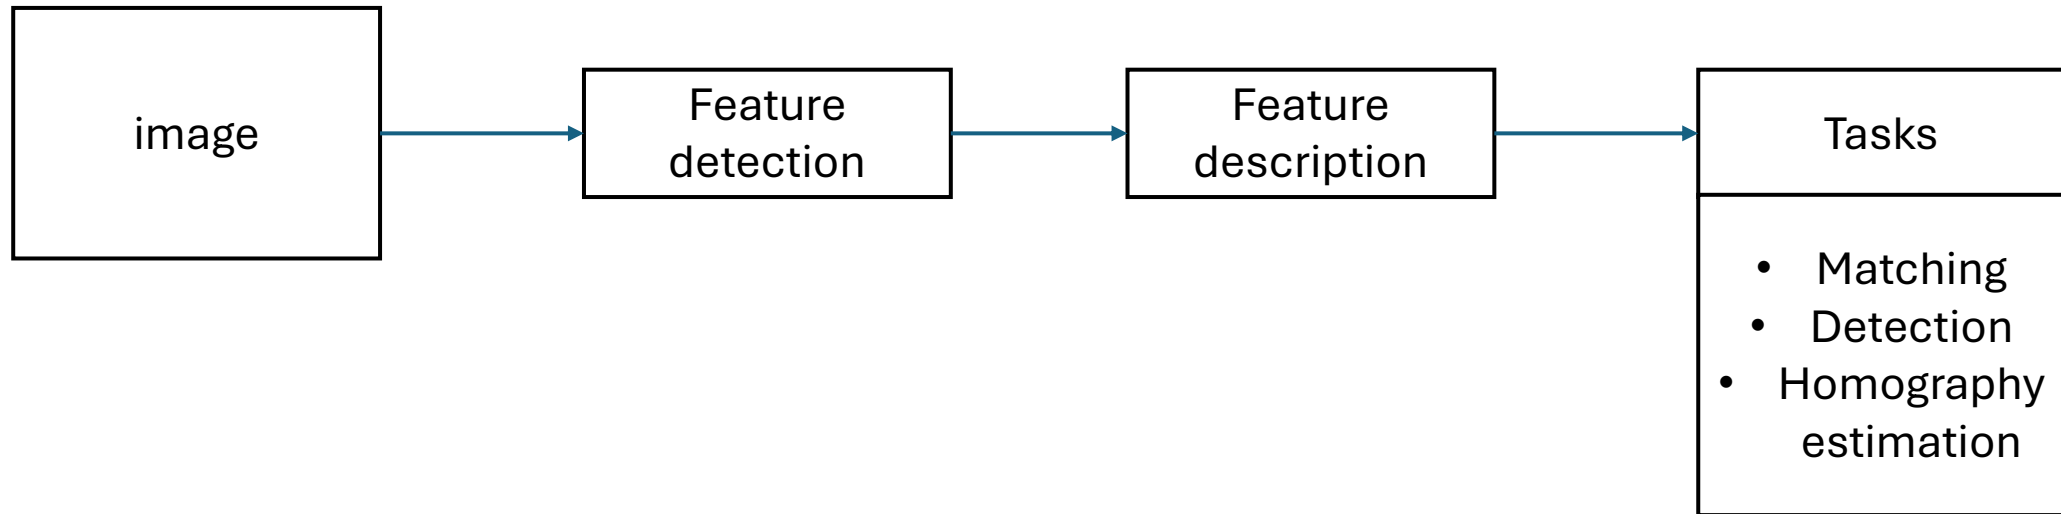
$f$  - Focal length  
 $z$  - depth  
 $c = (c_x, c_y)$  - optical center

$$Z \begin{bmatrix} u \\ v \\ 1 \end{bmatrix} = \begin{bmatrix} f_x & 0 & c_x \\ 0 & f_y & c_y \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} X \\ Y \\ Z \end{bmatrix}$$

# Pinhole camera model



# Classic understanding of images







# Object detection and tracking



# Matching for image stitching



# Homography estimation





# Edge detection


- Use derivatives (in x and y direction) to obtain pixels with high gradient
- Need smoothing to reduce noise prior to taking derivative
- E.g. Canny Edge Detector






# Corner Detection

- **Repeatability** – The same feature can be found in several images despite geometric and photometric transformations
- **Saliency** – Each feature is found at an “interesting” region of the image
- **Locality** – A feature occupies a “relatively small” area of the image




# Harris Corner Detector


Explore intensity changes within a window as the window changes location



"flat" region:  
no change in all  
directions




"edge":  
no change along the  
edge direction



"corner":  
significant change in  
all directions

# Harris Corner Detector – Code example OpenCV

```
import numpy as np
import cv2 as cv
filename = 'chessboard.png'
img = cv.imread(filename)
gray = cv.cvtColor(img,cv.COLOR_BGR2GRAY)
gray = np.float32(gray)
dst = cv.cornerHarris(gray,2,3,0.04)
#result is dilated for marking the corners, not important
dst = cv.dilate(dst,None)
# Threshold for an optimal value, it may vary depending on the image
img[dst>0.01*dst.max()]=[0,0,255]
cv.imshow('dst',img)
if cv.waitKey(0) & 0xff == 27:
    cv.destroyAllWindows()
```






# Difference of Gaussians (DoG)

Obtain Image in different scales and different amount of Gaussian blur

Keypoints obtained by computing difference of Gaussians in each scale

Choose best scale to represent that keypoint - the scale that contains a spatial gradient maxima



# Feature Descriptors


*Why do we need them?* To match relevant/corresponding image parts based on their feature similarity

## Properties

1. Information that is **invariant** w.r.t: illumination, pose, scale, intraclass variability
2. Highly **distinctive**: allows for finding the correct match with a good probability


# SIFT Descriptor

- Based on DoG keypoints
- Location and characteristic scale  $s$  given by DoG detector (**scale invariant**)



- Compute gradient at each pixel
- $N \times N$  spatial bins
- Compute an histogram  $h_i$  of  $M$  orientations for each bin  $i$

Obtain feature vector by concatenating all  $M \times N \times N$  orientations



- **Rotation invariant:** Find dominant orientation by building a orientation histogram. Rotate all orientations by the dominant orientation




# Other feature descriptors

- HoG (Histogram of oriented gradients)
- SURF (Speeded Up Robust Features)
- ORB (an efficient alternative to SIFT or SURF)
- FREAK (Fast Retina Keypoint)

# Summary of feature detectors/descriptors

- Based on spatial derivatives and local smoothing filters
- Based on expert knowledge
- Require some heuristic thresholds and filtering steps (not all keypoints are relevant)
- They are handcrafted to support desired properties, i.e. scale, rotation, illumination invariance.
- Hard to come up with rules that generalize well to all scenarios!
- What's next?

# Image classification performance



[Statistics provided by ILSVRC]




## ***Deep Learning***

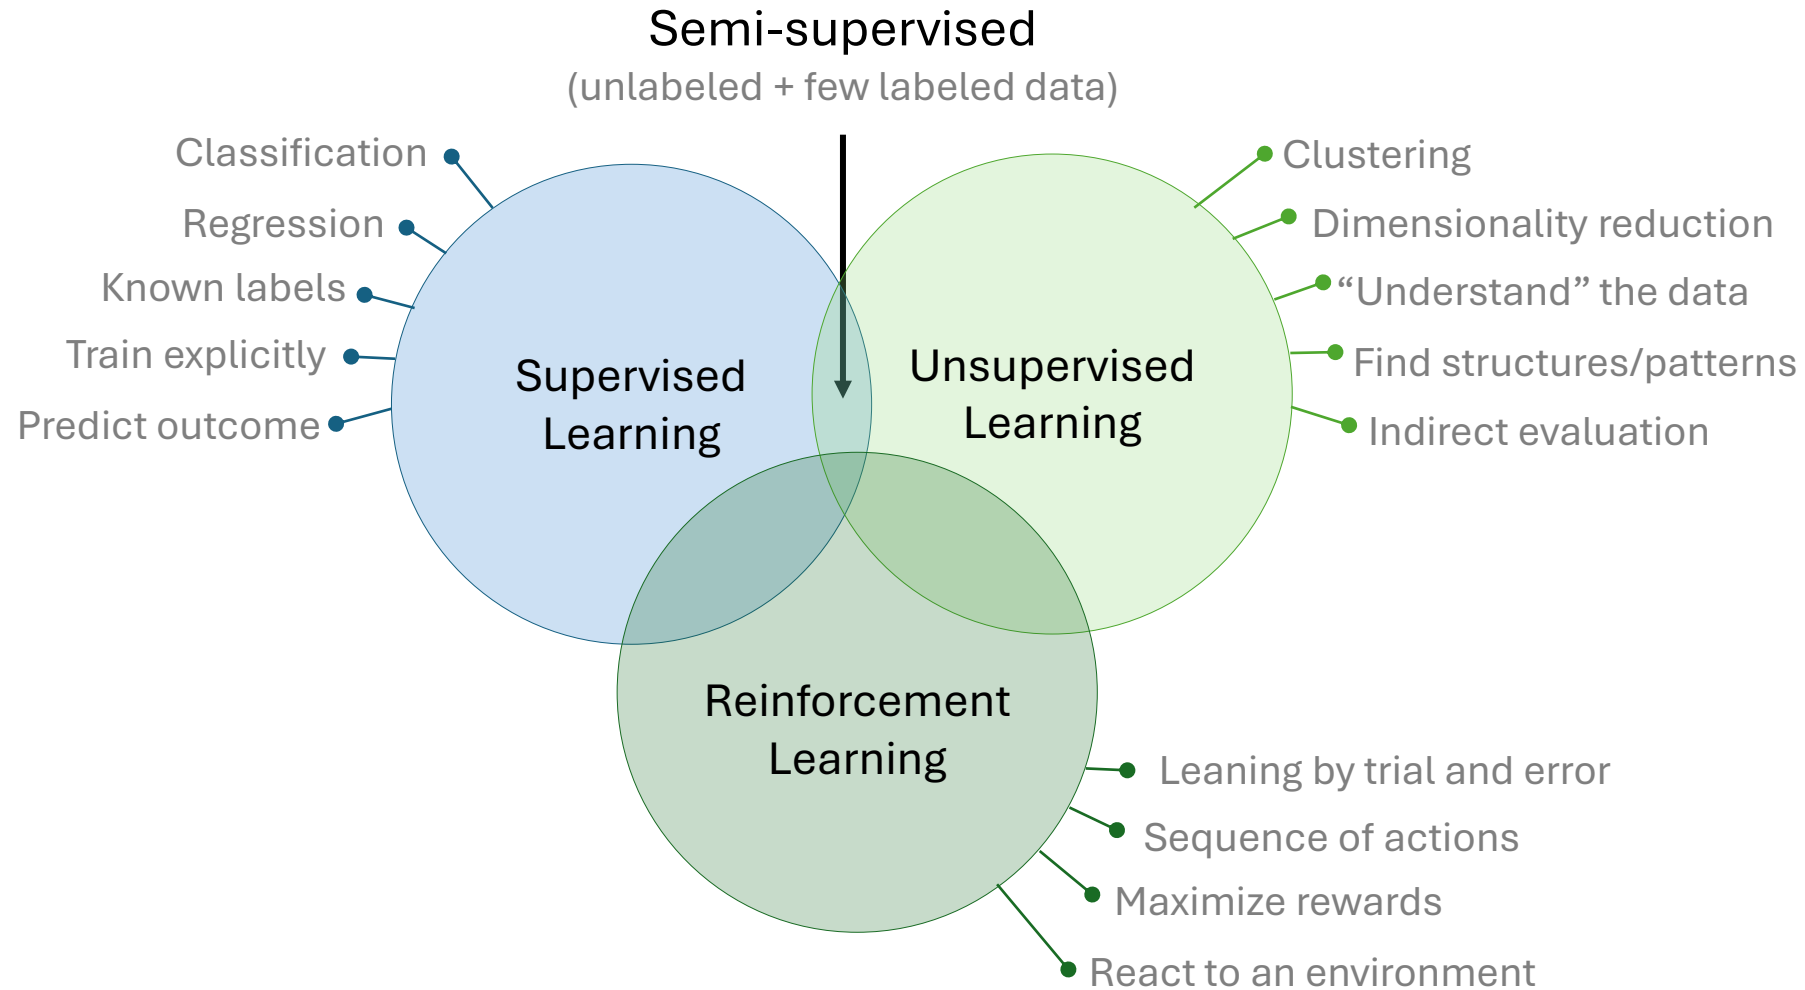
A type of machine learning based on artificial neural networks in which multiple layers of processing are used to extract progressively higher level features from data.

# Machine Learning

- [Arthur Samuel, 1959]
  - Field of study that gives computers the ability to learn without being explicitly programmed
- [Kevin Murphy]  
Algorithms that
  - automatically *detect patterns* in data
  - use the uncovered patterns to *predict* future data or other outcomes of interest
- [Tom Mitchell]  
Algorithms that
  - learn from experience (E)
  - with respect to some class of tasks (T)
  - to improve their performance (P)





# Machine Learning



# Supervised Learning

## Learning from Examples

- Training set of  $N$  samples  $(x^{(i)}, y^{(i)})$
- Generated by unknown function  $f$  s.t.  $f(x^{(i)}) = y^{(i)} \quad \forall i$
- $x^{(i)}$ : input,  $y^{(i)}$ : expected outcome
- Discover/Learn  $f^*$  that approximates  $f$
- Given a **new**  $x^{(j)}$  **with unknown**  $y^{(j)}$  compute  $y^{(j)}$  as  $y^{(j)} = f^*(x^{(j)})$




How can we assess the quality of the learned  $f$ ?

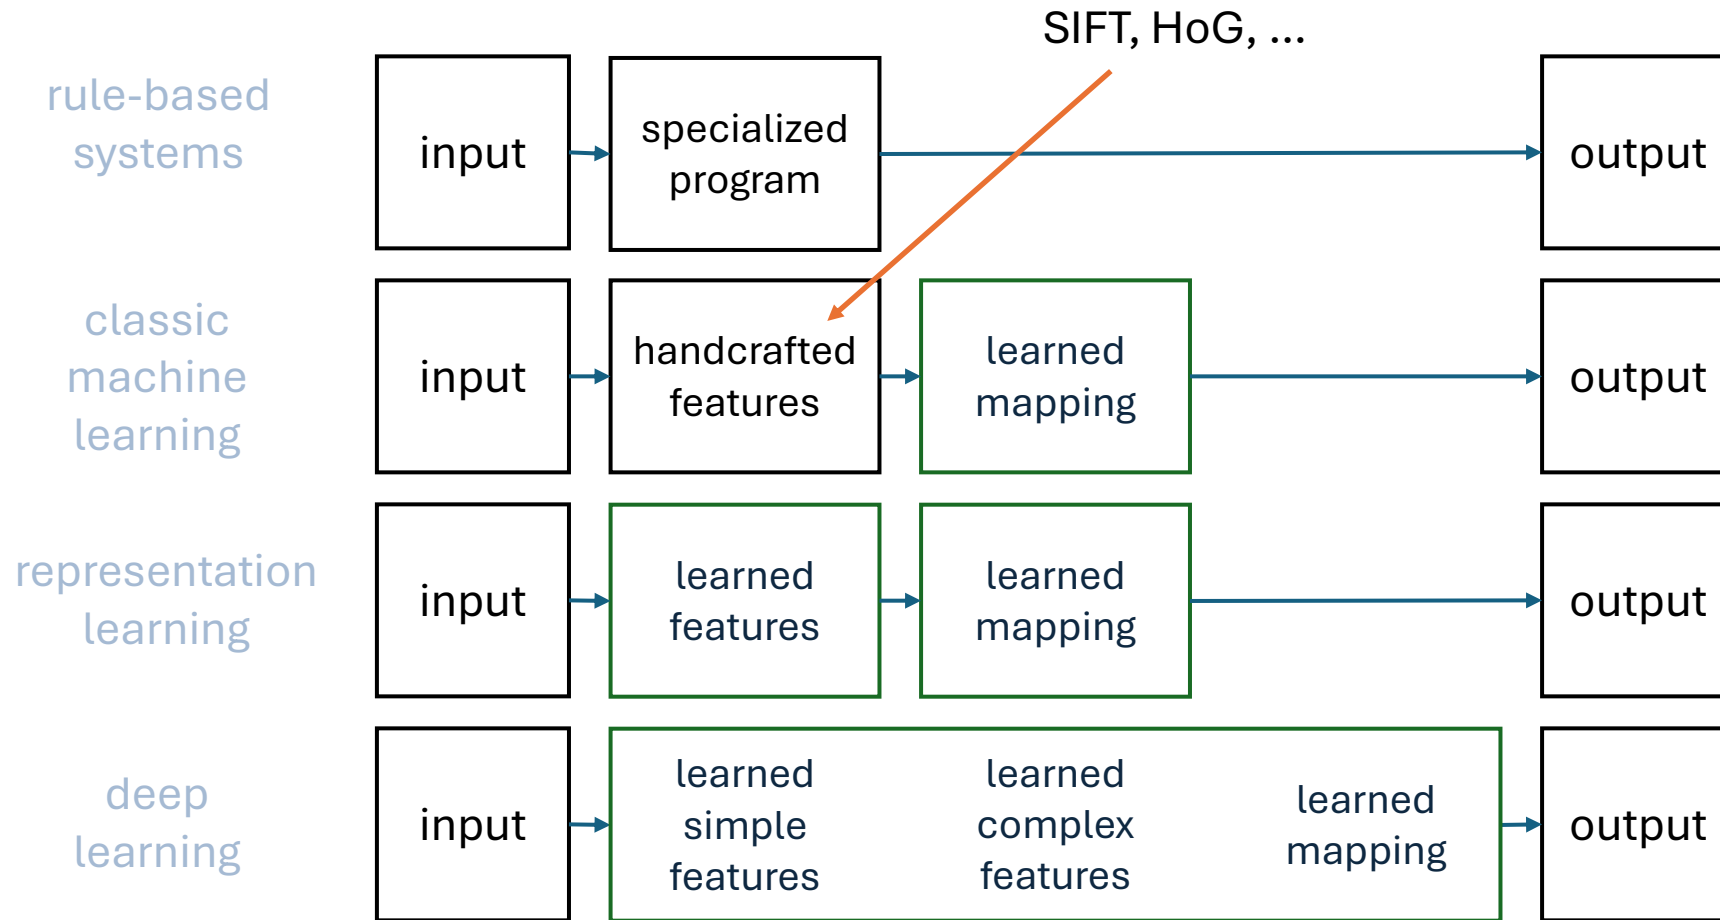


# Training – Validation – Testing

- Use only a subset of the samples for learning  $f$  (training set)
- The rest is for testing the quality of the predicted  $f^*$  (test set)
- If the learning process has parameters: split the training set again and tune the parameters on left-out subset (validation set)



# Why is deep learning attractive?



# Artificial Neural Networks

*"...a computing system made up of a number of simple, highly interconnected processing elements, which process information by their dynamic state response to external inputs."*

Dr. Robert Hecht-Nielsen in "Neural Network Primer: Part I" by Maureen Caudill, Feb. 1989


# The Perceptron

*“the embryo of an electronic computer that [the Navy] expects will be able to walk, talk, see, write, reproduce itself and be conscious of its existence.”*

Frank Rosenblatt acc. to New York Times from  
Mikel Olazaran (1996). "A Sociological Study of the Official History of the  
Perceptrons Controversy". *Social Studies of Science* 26 (3): 611–659

## Mark I Perceptron

- Frank Rosenblatt (1957)
- Image recognition
- 20 x 20 photo cells
- Learning with motors attached to potentiometers





# The Perceptron


$$f: \mathbb{R}^d \rightarrow \{1,0\}$$

$$f(x) = \begin{cases} 1 & \text{if } w \cdot x + b > 0 \\ 0 & \text{otherwise} \end{cases}$$


$$x, w \in \mathbb{R}^d \quad b \in \mathbb{R}$$

- Linear Classifier
- Only works well on linearly separable problem

# The Perceptron




# The Perceptron



To simplify notation:


- prepend a 1 to the input vector  $x$
- include the bias into the weights
- write everything as an inner product

# The Perceptron






# The Perceptron




What to do when the problem is not linearly separable?


# More Layers



# More Layers

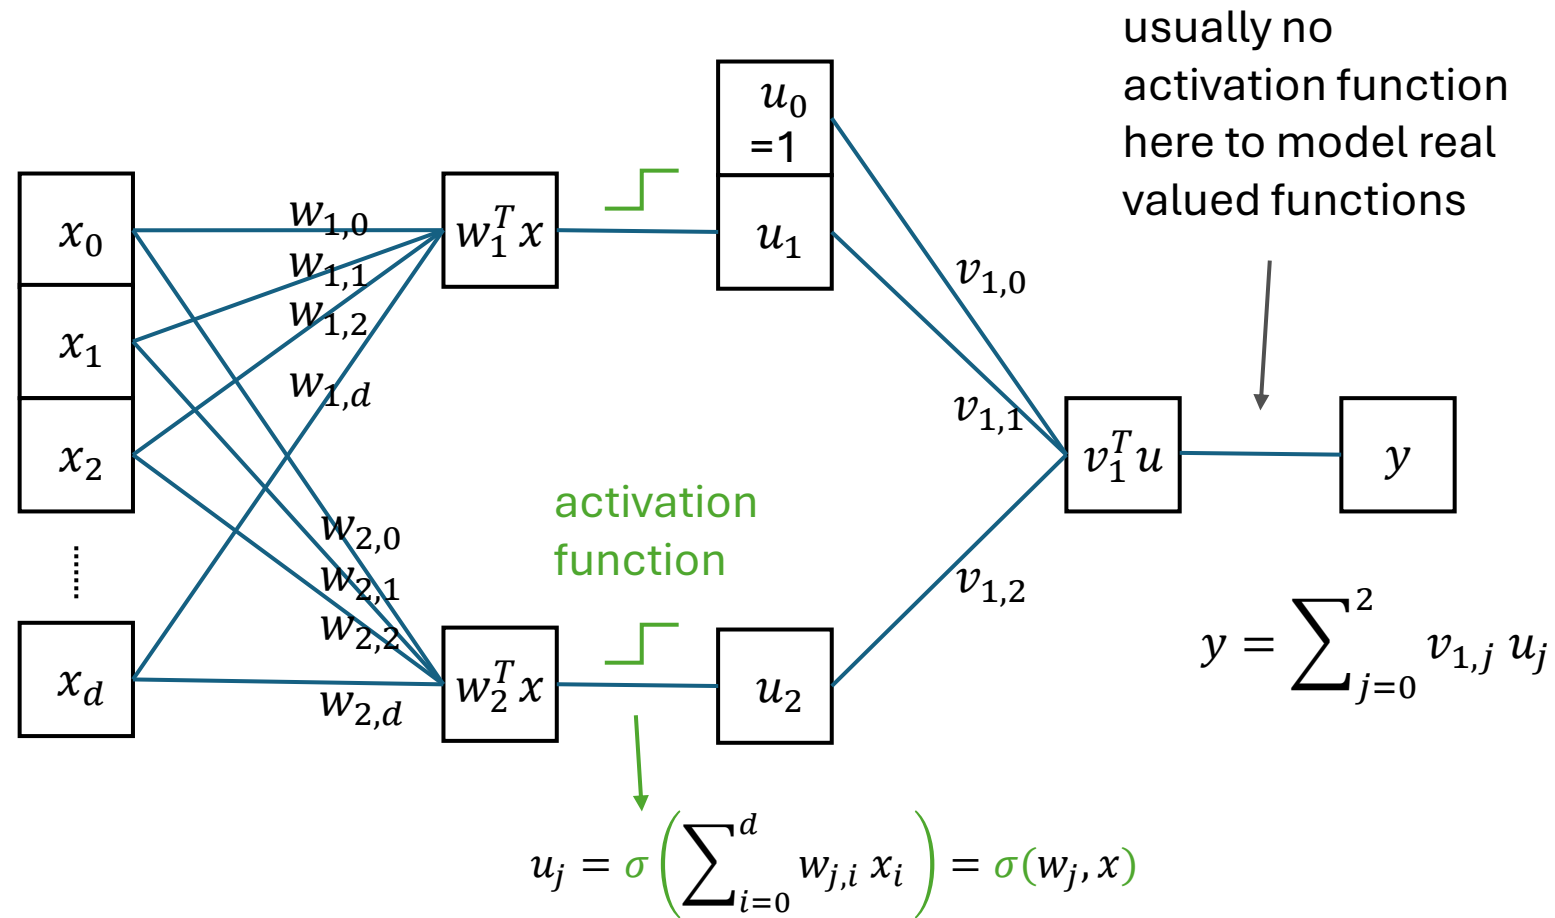


# More Layers





# More Layers



# More Layers

- Matrix notation greatly simplifies writing down the computations in the layers, e.g.

$$u_j = \sigma\left(\sum_{i=0}^d w_{j,i} x_i\right) = \sigma(w_j x)$$

$$u = \sigma(wx)$$

$$y = uv = v\sigma(wx)$$

- **Activation Function:** Do we need one?

Assume  $\sigma(z) = z$ :


$$y = \sigma(v\sigma(wx))$$
$$y = vwx = (vw)x$$

Yes, otherwise we still compute a linear function and nothing is gained by stacking the second layer!

# Activation function


Layers are typically followed by **non-linear** activation functions, that act per neuron.

**Sigmoid:**  $\sigma(x) = \frac{1}{1+e^{-x}}$



Output range: [0,1]

**Tanh:**  $\tanh(x) = 2\sigma(2x) - 1$




Output range: [-1,1]

# Activation function

Layers are typically followed by **non-linear** activation functions, that act per neuron.

## Rectified Linear Unit (ReLU)

$$f(x) = \max(0, x)$$




- Simply thresholds at zero
- Sparse activation
- Computationally efficient
- Non-saturating → speeds up convergence




# How to train your network

- Set of  $N$  samples  $(x^{(i)}, y^{(i)})$
- Define loss  $\mathcal{L}(f(x), y) \in \mathbb{R}$  to measure error *for a sample*
- *E.g. for regression task: mean square error (MSE) or mean absolute error (MAE) are common*
- **Training:** find weights that minimize  $\sum_i \mathcal{L}(f(x^{(i)}), y^{(i)})$  for the samples
- Often uses simple gradient descent methods
- **Backpropagation** computes the gradients of all parameters



# Gradient Descent

How to use **derivatives**?





- Gradient descent to minimize error function

$$w^{(t+1)} = w^{(t)} - \lambda \frac{\partial \mathcal{L}}{\partial w}$$

- Update the weights in every iteration of training with a small gradient step
  - Learning rate  $\lambda$  adjusts the step size
  - Error is defined over the **whole** training set:  $\sum_i \mathcal{L}(x^{(i)})$
  - Need to compute and sum the derivatives of all samples before *one* gradient step
  - Slow but accurate updates
- 
- Stochastic gradient descent (SGD)
    - Approximate derivative from small, random subset ([mini]batch) of training set
    - Noisy but faster
    - Usually: make sure to see every sample the same amount of times (epochs)

# Backpropagation

- Backpropagation is an efficient way to compute **derivatives** of  $\mathcal{L}$  w.r.t. *all* parameters
- Using (stored) **activations** from the forward pass
- Backpropagating information from layer  $i + 1$  to  $i$  and reusing already known (previously computed) derivatives
- Possible through chain rule



$$\begin{array}{l} \text{Layer 1} \\ \vdots \\ \text{Layer 3} \end{array} \quad \frac{\partial}{\partial x} f(g(h(x))) = \frac{\partial f(u)}{\partial u} (g(h(x))) \frac{\partial g(v)}{\partial v} (h(x)) \frac{\partial h(x)}{\partial x} (x)$$

$$\frac{\partial}{\partial u} f(g(h(x))) = \frac{\partial f(u)}{\partial u} (g(h(x)))$$

simpler (Leibnitz's notation)

$$t = f(u), \quad u = g(v), \quad v = h(x) \rightarrow \frac{\partial t}{\partial x} = \frac{\partial t}{\partial u} \frac{\partial u}{\partial v} \frac{\partial v}{\partial x}$$



Break 😊



# Deep Learning Frameworks

theano

MatConvNet



 PyTorch



[M]<sup>s</sup>  
MindSpore



Caffe2




TensorFlow



# PyTorch Overview

- **Open source** machine learning library based on the Torch library
- Operates on multi-dimensional vectors (**Tensors**)
- Can execute on the CPU, GPU, distributed systems, etc.
- **Dynamic** computational graph (can change on runtime)
  - Nodes: Tensors
  - Edges: mathematical operations




- Performs **automatic differentiation**
- Python and C++ interface
- **Torchvision**: package that implements many important vision algorithms

# PyTorch Overview a of Deep Learning Pipeline



Training loop



Inference

# PyTorch

`torch.nn.Module`

- Base class for a neural network module
- Can contain sub-modules
- Inherit this class when creating your neural network

`torch.nn.Parameter`

- Learnable tensor

`torch.nn.functional`

- A set of operations such as convolution, activation, etc.



# PyTorch Define a Neural Network Model

- Extend the `Module` class of `torch.nn`
- Implement the constructor and the forward member function

```
import torch.nn as nn

class ConvNet(nn.Module):
    def __init__(self):
        super(ConvNet, self).__init__()

        ...

    def forward(self, x):
        ...
```

- Optional: implement own `backward()` for custom back-propagation

# PyTorch Layers vs. Functions

## Layers

- Defined as **classes** in `torch.nn`
- Has attributes, like weights and bias
- Internally calls the functional API
- Use whenever possible, i.e. for standard layers
- In general good coding style

```
nn.Conv1d(in_channels, out_channels, kernel_size, stride=1, padding=0, dilation=1, groups=1, bias=True, padding_mode='zeros')
```

## Functions

- Defined as **functions** in `torch.nn.functional`
- **Only** provides **operation**, you need to pass your own weight and bias
- Learnable parameters need to be declared in `__init__()`, otherwise it will not learn
- Use in case you need to customize a layer

```
nn.functional.conv1d(input, weight, bias=None, stride=1, padding=0, dilation=1, groups=1)
```

# PyTorch Data Loader

Use python utilities: `torch.utils.data.DataLoader` class

- Represents a Python iterable over a dataset (pass a dataset as argument)
- Automatic batching in standard cases, use own `collate_fn()` to customize

```
from torch.utils.data import Dataset
from torch.utils.data import DataLoader
```

```
dataset = MyDataset()
dataloader = DataLoader(dataset, batch_size, ...)
```

```
class MyDataset(Dataset):
    def __init__(self, ...):
        ...

    def __getitem__(self, index):
        ...

    def __len__(self):
        ...
```

# PyTorch Data Loader

Use python utilities: `torch.utils.data.DataLoader` class

- Represents a Python iterable over a dataset (pass a dataset as argument)
- Automatic batching in standard cases, use own `collate_fn()` to customize

```
from torch.utils.data import Dataset
```

```
class MyDataset(Dataset):
```

```
    def __init__(self, data_dir):
```

```
        # get list of all image paths in the data_dir
```

```
        self.image_list = glob.glob(data_dir)
```

```
    def __getitem__(self, index):
```

```
        image, label = load_data(self.image_list[index])
```

```
        # normalization, augmentation, etc
```

```
        image, label = do_some_preprocessing(image, label)
```

```
        return image, label
```

```
    def __len__(self):
```

```
        return len(self.image_list)
```

# PyTorch Linear Regression

Given pairs of x and y data, learn w and b

Optimize using SGD  $y = wx + b$

Prepare toy data to train the model:

```
# in the dataset class
```

```
def __init__(self):
```

```
    # create toy data for training
```

```
    x_values = [0.1*i for i in range(100)]
```

```
    x_train = np.array(x_values, dtype=np.float32).reshape(-1, 1)
```

```
    y_values = [2*i + 1 + random.random()-0.5 for i in x_values]
```

```
    y_train = np.array(y_values, dtype=np.float32).reshape(-1, 1)
```


```
    # from numpy to torch tensors
```

```
    self.x_train = torch.from_numpy(x_train)
```

```
    self.y_train = torch.from_numpy(y_train)
```

```
def __getitem__(self, index):
```

```
    return self.x_train[index], self.y_train[index]
```



Load data



# PyTorch Linear Regression

We need the linear layer from **torch.nn**:

```
torch.nn.Linear(in_features, out_features, bias=True)
```

```
class LinearRegression(torch.nn.Module):  
    def __init__(self, inputSize, outputSize):  
        super(LinearRegression, self).__init__()  
        self.linear = torch.nn.Linear(inputSize, outputSize)  
  
    def forward(self, x):  
        out = self.linear(x)  
        return out
```

Define Model

# PyTorch Linear Regression

We need the linear layer from **torch.nn**:

`torch.nn.Linear(in_features, out_features, bias=True)`

Or create our own **w** and **b** parameters

```
class LinearRegression(torch.nn.Module):  
    def __init__(self, inputSize, outputSize):  
        super(LinearRegression, self).__init__()  
        # self.linear = torch.nn.Linear(inputSize, outputSize)  
        self.w = nn.Parameter(torch.ones([inputSize, outputSize]))  
  
        self.b = nn.Parameter(torch.zeros([outputSize]))  
  
    def forward(self, x):  
        # out = self.linear(x)  
  
        out = self.w * x + self.b  
        return out
```

Define Model

# PyTorch Linear Regression

```
#define data loader
dataset = MyDataset()
dataloader = DataLoader(dataset, batch_size=4)
# define model, loss function and optimizer
model = LinearRegression(inputSize=1, outputSize=1)
model.train()

loss_fn = torch.nn.MSELoss()
optimizer = optim.SGD(model.parameters(), lr=0.01)

# train loop
for i in range(n_epochs):
    for input, target in dataloader:
        optimizer.zero_grad()
        # forward step
        output = model(input)
        # compute loss
        loss = loss_fn(output, target)
        # optimize: compute gradients and apply
        loss.backward()
        optimizer.step()
# evaluate by printing w and b

print("slope: ", model.w.data.numpy(), "\t offset: ",
      model.b.data.numpy()) # slope: 1.9622    offset: 1.1127834
```


Forward step

Compute loss

Optimize

Evaluate

# How does the PyTorch model look like for this network?



```
class MyNet(torch.nn.Module):  
    def __init__(self):  
        super(MyNet, self).__init__()  
        ...  
  
    def forward(self, x):  
        out = ...  
        return out
```

# Google Colab



- A Jupiter notebook stored on Google drive
- It runs online on Google resources (no need to have your own GPU)
- We will use it for practical coding exercises
- Contains coding cells and text cells



# Coding exercise

Given a set of inputs  $x$  and a set of outputs  $y$ , learn the function  $y = f(x)$  by a neural network

Develop and experiment with 3 different network models and see how they compare

**Make a copy** of this Google Colab: <https://colab.research.google.com/drive/15wuKbpHuJmS8-FcsW2rRjuYmJh9aWIUb?usp=sharing>

Follow instructions and complete TODO list

Once you are done, share the link with me per email ([dl4cv.eci24@gmail.com](mailto:dl4cv.eci24@gmail.com)), using the **Share** button on the top right corner of Google Colab.

Run the code cells, and preferably let the running outputs there for me to see.

However, you can expect that I will try to run your coding cells myself to see if the output can be reproduced.

# References

Stanford CNN class notes:

<https://cs231n.github.io/>

General Machine Learning concepts explained simply:

<https://www.youtube.com/@statquest>

PyTorch tutorials:

<https://pytorch.org/tutorials/beginner/introyt.html>